

Section A: Subqueries (Scalar, Row, Column, IN, EXISTS, ANY, ALL)

1. Find the customers who have **never made a purchase** using `NOT IN` subquery.
 2. Display customer names who have made **purchases greater than the average sale amount**.
 3. Get the **total sales** for each customer using a correlated subquery.
 4. List customers who made **at least one purchase in 2022** using `EXISTS`.
 5. Retrieve customers who purchased **every available product** (using `ALL`).
 6. Show customers who have **more total purchases than customer ID 10** (using `ANY`).
 7. Get the name of the **top-spending customer** (using scalar subquery + `LIMIT`).
 8. Find customers from a city where **the average sale amount is above 600**.
 9. Show products that were **never purchased by customers in Chicago**.
 10. Display customers whose **sales exceeded their own average sale** (correlated subquery).
-

Section B: Views (Create, Replace, Drop, Updatable vs Read-Only)

11. Create a view to show only `customer_id`, `customer_name`, and `city`.
 12. Create a view `high_spenders` showing customers with total sales > 5000.
 13. Replace the `high_spenders` view to now include `city` as well.
 14. Drop the `high_spenders` view.
 15. Create a view to show all sales made in 2023 with product and sale amount.
 16. Create a read-only view using an aggregate function and explain why it's not updatable.
 17. Write a query using a view to show customer names and the count of purchases.
 18. Create a view that **hides customer join date** from analysts.
 19. Create a view for city-wise total sales, grouped by city.
 20. Demonstrate a role-based view: customers only from 'Los Angeles'.
-

Section C: Indexes (Create, Drop, Composite, Performance Impact)

21. Create an index on the `city` column of the `customers` table.
22. Create a composite index on (`customer_id`, `product`) in the `sales` table.
23. Drop the index on the `city` column.
24. Explain what happens when you insert 10,000 rows into a table with 5 indexes.
25. Create an index on `sale_date` and run a query to see performance difference with `EXPLAIN`.
26. Show an example where indexing a **low-cardinality column** (like `gender`) doesn't help.

27. Write a query where composite index on `(customer_id, product)` helps.
 28. Write a query where that same index doesn't help (e.g., only filtering by `product`).
 29. Create a query that benefits from both a `WHERE` and an `ORDER BY` index.
 30. Explain how too many indexes can slow down `UPDATE` or `DELETE` operations with examples.
-

Section D: Partitioning (Range, How to View, Insert, Query)

31. Create a `sales_partitioned` table partitioned by `YEAR(sale_date)` (range).
 32. Create at least **4 partitions** for years 2020–2023.
 33. Insert values into each partition manually and confirm they land in the correct one.
 34. Query `INFORMATION_SCHEMA.PARTITIONS` to list all partitions of the `sales_partitioned` table.
 35. Write a query to get **total sales in 2021** from partitioned table.
 36. Add a catch-all `pmax` partition for future sales.
 37. Explain how the **PRIMARY KEY must include partition key** in MySQL.
 38. Compare query performance for year-based filtering with and without partition.
 39. Show an example of how `DROP PARTITION` can be used to **archive old data**.
 40. Create a monthly partition for `sale_date` and verify inserts across months.
-

Section E: Integration + Higher-Order Thinking

41. Create a view that **combines partitioned sales data** and customer names via a `JOIN`.
42. Write a query using a **subquery inside a view** to show top 5 customers by sales.
43. Run `EXPLAIN` on a query with both a **view and an index** and analyze its efficiency.
44. Demonstrate how **index + partition + view** can be combined for a dashboard query.
45. Create a report query that uses: a view, a subquery, a join, and a filtered index (advanced full-scope test).